Introduction

The study of biota diversity has been a long-living topic fascinating a number of investigators in biological science ever since the time of natural history. Starting from recording the macrofauna and macroflora of an opening ecosystem, the focus on species recognition has been gradually shifted from the visible world to the micro ecosystem, the wonderland of microorganism. Unlike studying the species composition of macrobiota, the invisible nature of microbiota brings researchers a huge barrier in identifying and recording the composition of species diversity. The manageable microorganisms have once been the only group of species that can be identified from the standard protocol like, for example, isolating and culturing pure clones from a bunch of mixed species to determine their classification with biochemical analyses; those species that cannot be cultured then will leave no trace of themselves. Such workflows, unsurprisingly, depict only a cruel contour of the true picture of species diversity in microbiota; moreover, cell culturing is a time-consuming step, which seriously hampers the efficiency of workflow. Consequently, species identification without cell culturing became the critical point in build an efficient workflow in studying biodiversity of microbiota.

The emergence of metagenomics, however, solves the problem once and for all. Metagenomics refers to the study of retrieving and analyzing the genetic information from environmental samples directly without species identification in advance. A snapshot of species composition can now be recorded through the powerful sequencing technology and analyzed by a variety of bioinformatic tools. A series of software has been published and designed for different purposes. For example, PHACCS, a bioinformatic tool designed for estimating of the structure and diversity of uncultured viral communities; QIIME, an open-source bioinformatics pipeline for performing microbiome analysis including quality filtering, OTU picking, taxonomic assignment, phylogenetic reconstruction, and diversity analyses; MLST, a nucleotide sequence based approach for the unambiguous characterization of isolates of bacteria and other organisms via defining the distinct pattern of alleles combination in a series of loci involved in microbial metabolism. With the aid of metagenomics, the study microbiota has become a popular topics among a number of biological fields including environmental science, industrial fermentation, gastrointestinal health, and even mental disorders, showing the amazing linkage between the fluctuation of microbiota and a great deal of phenomena including soil fertility, diet preference, metabolomics, and human health. These studies unveil the hierarchical nature of the biological system, suggesting that the facts of macroscopic scale are tightly regulated by the composition of microbiota.

Whether the pipelines of metagenomics provide great help in documenting the biodiversity of multicellular communities would be a challenging and valuable trial. With the great success in the study of microbial diversity, metagenomics has been proved to have great potential in identifying biota diversity directly from environmental samples. Here, by using a RNA-seq dataset extracted from a

mixed collection of body fragments from several species of arthropod, we tested the efficiency of two pipelines, based on Basic Local Alignment Search Tool (BLAST) against all known sequences of NADH dehydrogenase subunit 6 from NCBI database, with either raw RNA seq or de novo transcripts assembly in advance, in identifying species composition of the meta-transcriptomic dataset. Collecting originally from an activity of BioBlitz, the popular field study of surveying and recording all the living species within a designated area, we have a complete inventory of the species contributing to the RNA-seq dataset as a standard in evaluating the efficiency of each of the pipelines.

More than showing the potential of metagenomics in species identifying biodiversity of arthropods, the workflows built in this study provide great value in agricultural application. Soil mesofauna, the study of tiny invertebrates such as mites, nematodes, or spiders that are usually 0.1 mm to 2 mm in length bears a great requirement for establishing the metagenomic pipelines for species identification directly from environmental samples. The complex communities of these small creatures, dwelling in soil or in a leaf cover layer on the soil surface, are hard to be completely covered and analyzed by traditional workflow based on morphology analyses, and, surprisingly, only a handful of publications have tried to address such questions with sequencing technology. Considering the great number of soil-transmitted pests that erode the total output value of agriculture globally, our work, contributing to species identification of invertebrates from meta-transcriptomic data, suggests a possibility of building a pipeline testing the presence of pests from soil samples. Our study shed light on applying metagenomics to species identification of invertebrates, providing a possible alternative of pest control in industrial agriculture.

Materials and Methods

RNA-seq Dataset

We use an RNA-seq dataset comprised of material collected from 13 insects and 2 outgroups (a spider and a millipede). It comprises 202,209,951, 100 bases-long, paired-end reads from a RNA-seq Illumina run.

Table 1: Above are the organisms reportedly sampled in this dataset. There are 13 insects and a spider and millipede inluded as outgroups. '-' represents a fields that is unknown.

Order	Family	Genus species	Common Name
Blattaria	Blattidae	-	cockroach
Coleoptera	Silphidae	Nicrophorus sp	carrion beetle
Diptera	Syrphidae	-	hover fly
Hemiptera	Reduviidae	Sinea diadema	spiny assassin bug

Order	Family	Genus species	Common Name
Hymenoptera	Formicidae	Camponotus pennsylvanicus	black carpenter ant queen
Hymenoptera	Vespidae	Polistinae Pollutes fuscatus	northern paper wasp male
Hymenoptera	Vespidae	Eumeninae Ancistrocerus sp.	mason wasp
Lepidoptera	Papilionidae	Papilio cresphontes	giant swallowtail butterfly
Trichoptera	Phryganeidae	-	caddisfly
Mantodea	Mantidae	Tenodera sinensis	Chinese mantis
Neuroptera	Chrysopidae	-	green lacewing
Odonata	Libellulidae	Sympetrum obstrusum	white-faced meadowhawk dragonfly
Orthoptera	Tettigoniidae	-	katydid
Arachnida	-	-	spider
Myriapoda	-	-	millipede

Mitochondrial dataset

We used the mitochondria currently published in NCBI RefSeq (ftp://ftp.ncbi.nlm.nih.gov/refseq/release/mitocho
There are 7652 species in total, and the number of species represented in each
clade in the table one is represented the Table 2.

Table 2: Column 1 (total) is a count of the number of species
in the NCBI common tree within the given clade. Column 2 is
the number of sequenced representatives. These numbers were
calculated using the `latide` script (`github.com/arendsee/latide`)
and `mitochondrial-stats.sh`.

clade	total	sequenced
Phryganeidae	65	1
Reduviidae	878	11
Trichoptera	10100	3
Syrphidae	1524	1
Formicidae	5272	13
Tenodera_sinensis	1	1
Silphidae	112	1
Myriapoda	1585	15
Tettigoniidae	990	30
Arthropoda	341615	1466
Arachnida	23163	127
Blattidae	108	3
Orthoptera	4337	126
Coleoptera	41773	93
Lepidoptera	93042	211
Papilionidae	721	7
Diptera	52124	145
Hymenoptera	59585	42

clade	total	sequenced
Vespidae	442	2
Mantodea	615	8
Mantidae	239	4
Neuroptera	1092	17
Chrysopidae	234	4
Hemiptera	16498	126

Quality check

The quality of the dataset was assessed using the software FastQC (Andrews, 2010) using the following command:

```
fastqc -t 2 -o /ptmp/660_project_bugs/fastqc_output \
    BioBlitz_NoIndex_L008_R1_001.fastq.gz \
    BioBlitz_NoIndex_L008_R2_001.fastq.gz
```

Pipeline I

First, we obtained the nucleotide sequence for every mitochondrial genome present into NCBI RefSeq database from (ftp://ftp.ncbi.nlm.nih.gov/refseq/release/mitochondrion/, specifically the files "mitochondrion.1.1.genomic.fna.gz" and "mitochondrion.2.1.genomic.fna.gz"). We then aligned the dataset against this mitochondrial genomic database using Bowtie2 (Langmead and Salzberg, 2012). For this we used the following commands:

For concatenate both files

```
cat mitochondrion.*.genomic.fna > mitochondrion.all.fna
```

For building the bowtie2 reference genome index

```
bowtie2-build mitochondrion.all.fna mito_gen
```

For making the alignment

```
bowtie2                                     \
    -p 32                                   \
    --al bugs_unp_al.fastq                  \
    --al-conc bugs_conc_al.fastq            \
    --un bugs_unp_unal.fastq                \
    --un-conc bugs_nonconc.fastq            \
```

```
-x mitoGen/mito_gen                      \
-1 BioBlitz_NoIndex_L008_R1_001.fastq.gz \
-2 BioBlitz_NoIndex_L008_R2_001.fastq.gz \
-S bugs_vs_mito.sam
```

The –al and –al-conc arguments output the reads that align unpaired and concordant, respectively, to the database. Whereas the –un and –un-conc arguments output the reads that have one unaligned pair or both pairs unaligned, respectively.

We then performed a "de novo" transcript assembly on the concordantly aligned reads obtained from the Bowtie2 alignment step. This was achieved using Trinity (Grabherr et al, 2011) with the following command:

```
Trinity                                                \
    --CPU 48                                           \
    --max_memory 10G                                   \
    --seqType fq                                       \
    --left /ptmp/660_project_bugs/bugs_conc_al.1.fastq \
    --right /ptmp/660_project_bugs/bugs_conc_al.2.fastq
```

For the next step, we extracted the aminoacidic sequence of the mitochondrial protein "NADH dehydrogenase subunit 6" (ND6) for every specie in NCBI RefSeq database (ftp://ftp.ncbi.nlm.nih.gov/refseq/release/mitochondrion/, specifically from mitochondrion.1.protein.faa.gz and mitochondrion.2.protein.faa.gz). We then performed a tBLASTn alignment on our previously "de novo" assembled transcripts versus the ND6 proteins using BLAST+ software (Camacho et al. 2009). For this we used the following commands:

For making the BLAST database

```
makeblastdb                              \
    -in trinity_out_dir/Trinity.fasta    \
    -dbtype nucl                         \
    -out blast_trinityDB                 \
    -title trinity_out
```

For making the tBLASTn alignment

```
tblastn                                                   \
    -task tblastn-fast                                    \
    -query nd6_2-mito.faa                                 \
    -db /ptmp/BCB660/660_project_bugs/blast_trinityDB     \
    -evalue 0.001                                         \
    -num_threads 16                                       \
```

```
    -max_target_seqs 1000000000                           \
    -outfmt '7 qseqid sseqid qlen slen
             qstart qstop sstart sstop
             bitscore pident positive evalue' \
    > nd6_2-blast.tab
```

After getting the BLAST results, we used a bash script to re-format the obtained
file for an easier extraction of the required data. This script will remove the
comment lines from the BLAST output and will add a column with the name of
the specie for every ND6 protein accession number (see supplementary materials
for scripts).

A summary of the script is provided below (for a full detailed and commented
version, see supplementary materials):

```
join -t $'\t' \
    <( cat $protein_reference               |
       sed -nr 's/>([^ ]+).*\[(.*)\]/\1\t\2/p' |
       tr ' ' '_' | sort ) \
    <( grep -Pv '^#' $input_blast_result | sort ) \
    > $output_blast_result
```

Once obtained the re-formatted BLAST results, we calculated the most repre-
sented species in the BioBlitz sample using a R script (see supplementary materi-
als for the commented script). The resulting file is a list of the most represented
species, which was used to generate a phylogenetic tree using NCBI Common Tree
web site (https://www.ncbi.nlm.nih.gov/Taxonomy/CommonTree/wwwcmt.cgi).
Finally, the obtained tree was visualized using TreeGraph2 software (Stöver and
Müller, 2010).

Pipeline II

For this, we first used the provided reads in the dataset to build a BLAST
database using the following command (for the full script, please refer to supple-
mentary materials):

For converting the reads from fastq to fasta format

```
zcat $r1 $r2 |
    awk 'NR % 4 == 1 {print ">" $0} NR % 4 == 2 {print}' > $tmpfa
```

For building the BLAST database

```
makeblastdb                    \
    -in $tmpfa                 \
    -dbtype nucl               \
    -out blastdb/bioblitz \
    -title bioblitz
```

After building the BLAST database we aligned the mitochondrial protein "NADH dehydrogenase subunit 6" (ND6) dataset (the same dataset from pipeline 1) to it using the following command:

```
tblastn                                                        \
   -task tblastn-fast                                          \
   -query nd6_2-mito.faa                                       \
   -db /ptmp/BCB660/arendsee/bcb660-project/blastdb/bioblitz \
   -evalue 0.001                                               \
   -num_threads 16                                             \
   -max_target_seqs 1000000000                                 \
   -outfmt '7 qseqid sseqid qlen slen qstart
            qstop sstart sstop bitscore pident
            positive evalue' \
   > nd6_2-blast-all.tab
```

After getting the BLAST results, we performed the same steps as in Pipeline I. Briefly, used a bash script to re-format the obtained file, calculated the most represented species in the BioBlitz sample using a R script, generate a phylogenetic tree using NCBI Common Tree web site and visualized the tree using TreeGraph2 (as stated before, for a full description of the scripts please refer to the supplementary materials).

Results / Discussion

Quality check: The FastQC run reported the following errors: Per base sequence quality (there was a considerable drop in base call quality at the last bases), Per base sequence content, Sequence duplication levels and Kmer content (figure 1). Since our project is to identify the different species present in a BioBlitz sampling, we decided to keep all reads without performing any trimming or filtering process on them, regardless of the quality check results.

Pipeline I: From the downloaded RefSeq mitochondrion database we used all of 7,809 available genome sequences to build our Bowtie2 reference index. The Bowtie2 alignment gave us a total of 8,949,517 paired reads that aligned concordant to our mitochondrial genomes database (roughly 4.4% of the total reads). The Trinity "de novo" assembly on these files predicted 534 transcripts. These transcripts were used as the database for our BLAST search which gave us a
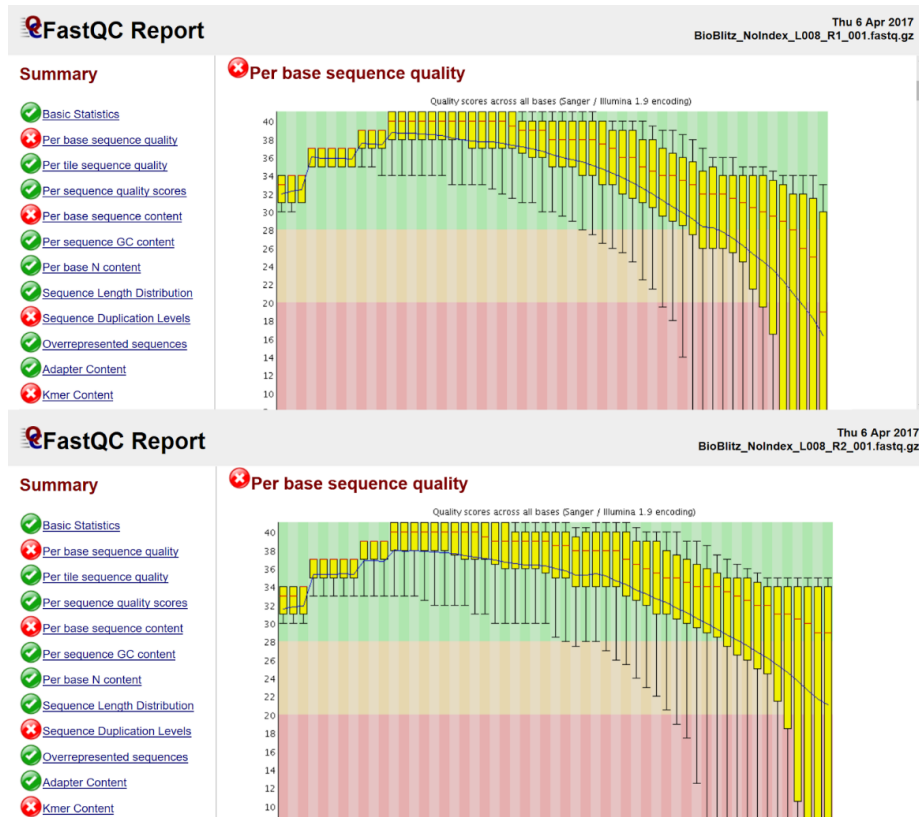
Figure 1: Results from FastQC analysis performed on the provided dataset. On the left side of each panel we can appreciate the result for each test, failed tests are depicted by red circles with a white cross. On the right side of each panel we see the phred value (average and median) per nucleotide position. Upper panel: analysis on forward reads; lower panel: analysis on reverse reads.

8

total of 7,486 hits. Using tBLASTn alignment score and results filtering, we could indicate how many times a certain specie's ND6 protein had a maximum score hit for any different predicted transcript. We used these results to suggest this specie's RNA abundance in out BioBlitz sample. From this data, we obtain that 12 species were likely to be present. Every specie was represented either 1 or 2 times at most. The exception was one specie, represented 25 times (the butterfly Tirumala limniace, table 1 and figure 2).

Table 3: List of species present according to results from Pipeline I

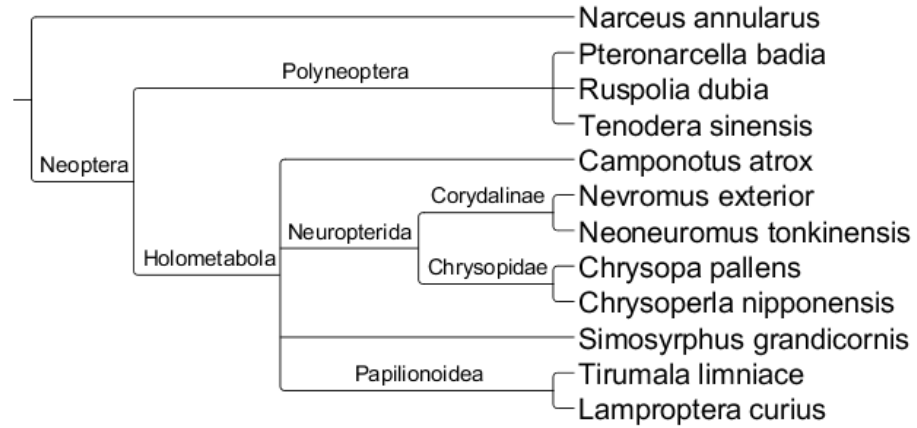| Specie Name	Max-score hits count
Camponotus atrox	2
Chrysopa pallens	1
Chrysoperla nipponensis	1
Lamproptera curius	1
Narceus annularus	1
Neoneuromus tonkinensis	1
Nevromus exterior	1
Pteronarcella badia	1
Ruspolia dubia	1
Simosyrphus grandicornis	2
Tenodera sinensis	1
Tirumala limniace	25

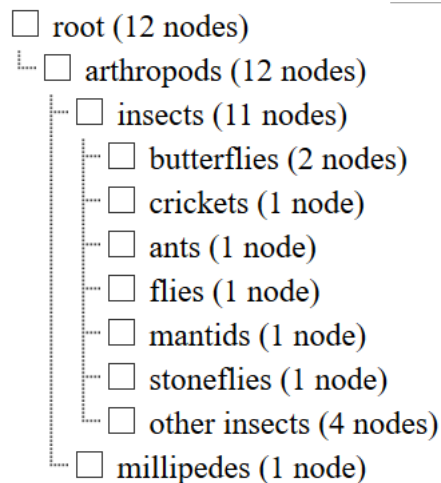Figure 2: Phylogenetic tree obtained from Pipeline I

Figure 3: different sub-categories of species present in Pipeline I, according to NCBI taxonomy web site

Pipeline II

We performed a tBLASTn using the entire 202,209,951 BioBlitz dataset read pairs as our database and all of the 7,418 ND6 protein sequences available on NCBI RefSeq as our query. The BLAST alignment produced 4,646,398 total hits. Again, using tBLASTn alignment score and results filtering, we could indicate how many times a certain specie's ND6 protein had a maximum score hit for any different read of our BioBlitz dataset. We used these results to suggest this specie's RNA abundance in our sample. From this processing, we obtained a total of 214 species with a maximum-score hit. In order to narrow down our results, we applied an extra filter for extracting only those species with a number of maximum hits above a defined threshold. Table 2 shows the different selected threshold values and the number of species obtained. We could observe that setting our max-score hits threshold above 30 showed little effect on the number of species extracted. Nevertheless, we decided to use 200 max-score hits as threshold value to narrow as much as we can our output. Figure 3 shows the obtained tree.

Table 4: Number of species showing different max-score hit thresholds

Number of max-score hit threshold	Number of species
0	214
10	78
20	61
30	52

Number of max-score hit threshold	Number of species
40	49
50	49
60	47
70	46
100	40
200	33

Figure 3. Phylogenetic tree obtained from Pipeline II (Right) and the different sub-categories present, according to NCBI taxonomy web site (Lower-left).
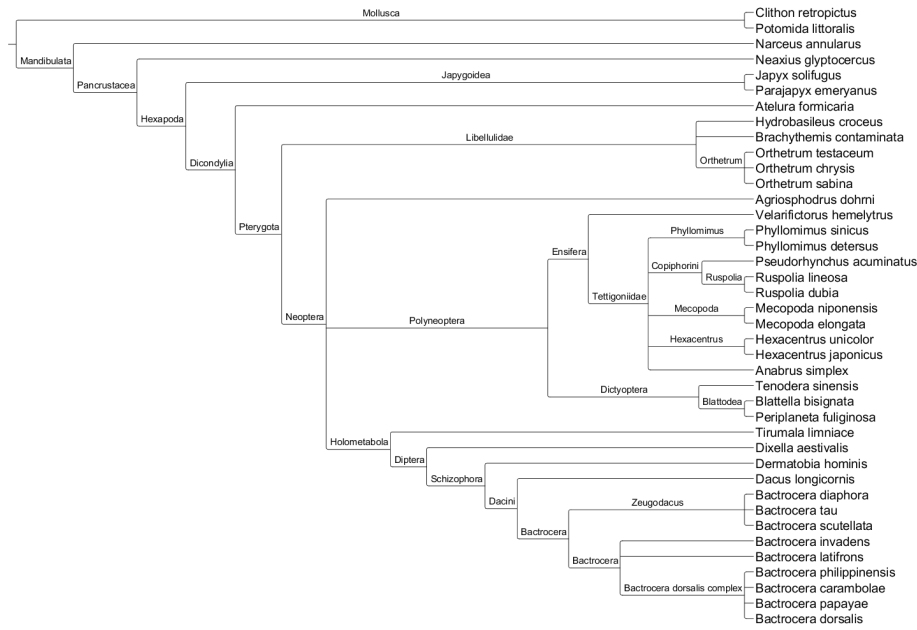

Figure 4: Phylogenetic tree obtained from Pipeline II

Comparison to expected species

All analysis to this point has assumed nothing was known about the sample that was sequenced. This, of course, is not true, since the researchers sampled specific organisms (see Table 1). Of the species listed in Table 1, only *Tenodera sinensis* has a mitochondria that has been sequenced. As expected, this species is identified exactly by both pipelines. A sister species of *Componatus pennsylvanica*, *Componatus atrox*, is found.

- □ root (40 nodes)
  - □ animals (40 nodes)
    - □ arthropods (38 nodes)
      - □ insects (36 nodes)
        - □ flies (12 nodes)
        - □ crickets (11 nodes)
        - □ dragonflies (5 nodes)
        - □ roaches (2 nodes)
        - □ butterflies (1 node)
        - □ mantids (1 node)
        - □ bugs (1 node)
        - □ other insects (3 nodes)
      - □ crustaceans (1 node)
      - □ millipedes (1 node)
    - □ molluscs (2 nodes)
      - □ gastropods (1 node)
      - □ bivalves (1 node)

Figure 5: different sub-categories of species present in Pipeline II, according to NCBI taxonomy web site

Conclusions

Here, we assessed two different methods for stablishing the number and type of species in a BioBlitz sampling. The first method relied on a more conservative pipeline, using "flagship" next generation sequencing analysis software such as Bowtie2 and Trinity. The second method involved a very simplified, yet elegant, pipeline using BLAST and locally-made scripts (bash and R scripts). Each method provided enough data to suggest a list of species present however we can appreciate the following differences/limitations: while the pipeline I gave us much less species (we only obtain 12 species compared to the 214 obtained with pipeline II) it seems to be very accurate on the number and type of species. On the other hand, our pipeline II needed an extra filtering step. However, having a high number of species on pipeline II can give us much more flexibility for performing further analysis that we may want (such as determining non-insect species presence our sampling). As for required time and resources, both pipelines performed very similar, running the entire pipeline in about 2 days on the hpc cluster.

From these results, we only can conclude that both pipelines are useful for identifying the present species on a BioBlitz sampling. Whether you choose to use one or the other pipeline may depend on software availability and the research question, among other issues related to the experiment being conducted.

Author's Contribution

Zebulun Arendsee wrote the scripts for taxonomic analysis and blast. Ang-Yu Liu performed BUSCO analyses (though we did not include these results, they represented a great deal of work) and also performed exhaustive literature review. Christian Montes performed the fastqc, trinity, bowtie analyses, and performed ran and analyzed scripts form ZA. All of the authors contributed to the design and writing of this project.

References

Andrews S. (2010). FastQC: a quality control tool for high throughput sequence data. Available online at: http://www.bioinformatics.babraham.ac.uk/projects/fastqc

Camacho, C., Coulouris, G., Avagyan, V., Ma, N., Papadopoulos, J., Bealer, K. and Madden, T.L., 2009. BLAST+: architecture and applications. BMC bioinformatics, 10(1), p.421.

Grabherr, M.G., Haas, B.J., Yassour, M., Levin, J.Z., Thompson, D.A., Amit, I., Adiconis, X., Fan, L., Raychowdhury, R., Zeng, Q. and Chen, Z., 2011. Trinity:

reconstructing a full-length transcriptome without a genome from RNA-Seq data. Nature biotechnology, 29(7), p.644.

Langmead, B. and Salzberg, S.L., 2012. Fast gapped-read alignment with Bowtie 2. Nature methods, 9(4), pp.357-359.

Stöver, B.C. and Müller, K.F., 2010. TreeGraph 2: combining and visualizing evidence from different phylogenetic analyses. BMC bioinformatics, 11(1), p.7.

Scripts

Most scripts are available on `github.com/arendsee/bcb660-project`. These scripts were wrapped for slurm submission with the `autoslurm.sh` script (not included, but on github).

Scripts used:

1. Get mitochondrial genomes

```
# Retrieve mitochondrial genomes
# If they are already downloaded, do nothing

outdir=data_mitochondria

if    [[ ! -d $outdir ]]
then
    mkdir $outdir
    wget \
    -P $outdir -nd -r -A gz \
    ftp://ftp.ncbi.nlm.nih.gov/refseq/release/mitochondrion
    cd $outdir
    gunzip *gz
fi
```

2. Build a BLAST database

```
module load ncbi-blast

[[ -d blastdb/ ]] || mkdir blastdb

r1=/ptmp/bioblitz/BioBlitz_NoIndex_L008_R1_001.fastq.gz
r2=/ptmp/bioblitz/BioBlitz_NoIndex_L008_R2_001.fastq.gz

# Check existence of the input files, die if missing
[[ -f $r1 && -f $r2 ]] || ( echo "Cannot open fastq files: $r1 $r2" >&2 && exit 1 )
```

```
tmpfa=fastq-fasta-extract_DELETEME.fa

# If z.fa already exists, and isn't an empty file, don't remake it
if [[ ! -s $tmpfa ]]
    then
    # Extract and combine the fasta sequences from the two fastq files
    # NOTE: I am not doing any cleaning or trimming here
    zcat $r1 $r2 |
    awk  'NR % 4 == 1 {print ">" $0} NR % 4 == 2 {print}' > $tmpfa
fi


if [[ ! -d blastdb ]]
then
    # Make the blast database, to access it with blast, use a command of the form:
    # $ blastp -db /blastdb/bioblitz -query whatever.faa [options]
    makeblastdb -in $tmpfa -dbtype nucl -out blastdb/bioblitz -title bioblitz
    echo "making a db" > &2
    else
    echo "not making a db, already exists" > &2
fi

# If makeblastdb fails,
if [[ $? -eq 0 ]]
    then
    rm $tmpfa
    exit 0
else
    echo "makeblastdb run failed" >&2
    exit 1
fi
```

   3. run the tBLASTn alignment

```
module load python/3.6.0
module load ncbi-blast

[[ -d smof ]] || git clone https://github.com/arendsee/smof

all_pro='data_mitochondria/mitochondrion*protein.faa'

base=nd6_2
protein_name="NADH dehydrogenase subunit 6";

blastresult=${base}-blast.tab
```

```
my_pro=data_mitochondria/${base}-mito.faa

cat $all_pro | smof/smof.py grep "$protein_name" > $my_pro

tblastn                          \
-task tblastn-fast               \
-query $my_pro                   \
-db blastdb/bioblitz             \
-evalue 0.001                    \
-num_threads 8                   \
-max_target_seqs 1000000000 \
-outfmt '7 qseqid sseqid qlen slen qstart qstop sstart sstop bitscore pident positive evalue
```

4. Re-formatting BLAST output (getting species names)

```
#!/usr/bin/bash

set -u

# USAGE:
#    # assuming you are running this from the bcb660-project folder:
#    scripts/autoslurm.sh scripts/add-names-to-blast-result.sh nd6-blast.tab

# This should be the name of the output file from BLAST, e.g. nd6-blast.tab
# 1. The file may have comments inside
# 2. The file is expected to have the extension `.tab`
input_blast_result=$1

# --- You shouldn't need to change this
# The protein fasta files are used just to extract the map between reference
# and species name. The fasta headers are assumed to have the format:
# >REFERENCE ... [ SPECIES_NAME ]
protein_reference=data_mitochondria/mitochond*protein.faa

# This is the name of the output file, the parameter expansion below replaces
# the extension on the input.
output_blast_result=${input_blast_result/.tab/-with-names.tab}

# Extract a reference to name map from the protein sequence file
# Join the map table with the blast result table
#  1. This adds the scientific species names as the second column in the output file
#  2. Replaces space in the species name with underscores
#  3. This will remove all comments from the blast result file
```

```
join -t $'\t' \
<( cat $protein_reference                    |
sed -nr 's/>([^ ]+).*\[(.*)\]/\1\t\2/p' |
tr ' ' '_' | sort ) \
<( grep -Pv '^#' $input_blast_result | sort ) > $output_blast_result
```

5. R script for extracting the most represented species

```r
require(readr)
require(dplyr)
require(magrittr)

# Provide the name for the re-formatted BLAST output
blastfile <- 'nd6-blast-with-names.tab'

d <- read_tsv(blastfile, col_names=FALSE, comment="#") %>%
select(X2,X3,X8) %>%
rename(species=X2,read=X3,bitscore=X8)
d$species <- gsub("_", " ", d$species)

# total number of species
(total <- length(unique(d$species)))

# species list (you can plug this into NCBI common tree)
write(unique(d$species), file="species-with-hits.txt")


# Get just the matches that are maximum for each read
d <- group_by(d, read) %>% dplyr::filter(bitscore == max(bitscore))

# number of species with maximum against at least one read
(withmax <- length(unique(d$species)))

(counts <- summary(factor(d$species), maxsum=Inf))
write(unique(d$species), file="species-with-a-max-hit.txt")

hist(log(counts))

# modify this to get the list of species with more than N max-score hits
names(counts[counts > 200]) %>%
write(file="species-with-more-than-200-max-hits.txt")
```

6. `mitochondrial-stats.sh` - Clade membership analysis

```
set -u
```

17

```
mitodir=data_mitochondria
sample=sample-taxonomies.tab

statdir=mito-stats

[ -d $statdir ] || mkdir $statdir

species_list=$statdir/species-list.txt
gene_counts=$statdir/gene-counts.tab
irregular_names=$statdir/irregular-names.tab
mitorep=$statdir/mitochondria-report.txt

mitoclades=$statdir/mito-clade-counts.tab

mitoprot=$mitodir/*protein*faa

list-species () {
    sed -nr 's/^>.*\[(.*)\]$/\1/p' $mitoprot |
        uniq | sort -u
}

count-genes () {
    echo -e "count\tn_proteins"
    sed -nr 's/^>.*\[(.*)\]$/\1/p' $mitoprot |
        sort | uniq -c |
        awk '{print $1}' |
        sort -n | uniq -c | sort -rn |
        sed 's/^ *//; s/  */\t/'
}

list-irregular-species-names () {
    grep -vP '^[A-Z][a-z]+ [a-z]+$' $species_list | sed 's/^/  /'
}

make-report () {
    echo -n "Number of represented mitochondrial genomes: "
    wc -l $species_list | sed 's/ .*//'

    echo
    echo "First 10 most common mitochondrial gene counts"
    head -11 $gene_counts

    echo
    echo "Species with irregular names (i.e. not '[A-Z][a-z]+ [a-z]+'):"
    wc -l $irregular_names | sed 's/ .*//'
```

```
}

[ -f $species_list    ] || list-species                    > $species_list
[ -f $irregular_names ] || list-irregular-species-names > $irregular_names
[ -f $gene_counts     ] || count-genes                     > $gene_counts
[ -f $mitorep         ] || make-report                     > $mitorep



# ==================== T A X O N O M I C   R E P O R T =======================

if [ ! -d litade ]
then
    git clone https://github.com/arendsee/litade
    cd litade
    ./setup.sh
    cd ..
fi
litade=$PWD/litade/litade.pl

if [ ! -f taxid2sciname.tab ]
then
    wget -O taxdmp.zip ftp://ftp.ncbi.nih.gov/pub/taxonomy/taxdmp.zip
    unzip -p taxdmp.zip names.dmp > names.dmp
    # Build a map from Taxon Id to Scientific name
    sed 's/[\t ]*|[\t ]*/\t/g' names.dmp |
        awk '
            BEGIN{FS="\t"; OFS="\t"}
            /scientific name/ {print $1,$2}
        ' > taxid2sciname.tab
fi

sample_taxids="inputs/sample-taxids.tab"
mito_taxids=mito-taxids.txt
protein_gpff=$mitodir/*protein*gpff

join -t $'\t' \
     <(sort $sample_taxids) \
     <($litade $(cut -f 1 $sample_taxids) | sort) |
    tee z |
    awk 'BEGIN{FS="\t"; OFS="\t"} {print $1,$2}' | sort | uniq -c |
    sed 's/  *//' | sed 's/ /\t/' |
    awk 'BEGIN{FS="\t"; OFS="\t"} {print $2,$3,$1}' | sort > a

join -1 1 -2 3 -t $'\t' \
     <(sed -nr 's/.*db_xref="taxon:([0-9]+).*/\1/p' $protein_gpff | sort -u) \
```

19

```
   <(sort -t $'\t' -k3 z) |
   cut -f2 | sort | uniq -c |
   awk 'BEGIN{OFS="\t"} {print $2,$1}' > b


echo -e "taxid\tclade\ttotal\tsequenced" > $statdir/taxid-counts.tab
join -t $'\t' <(sort a) <(sort b) >> $statdir/taxid-counts.tab

rm a b z
```